Game	Key parameters	Utility
trade_comm_n4	num_items=4	General-sum
goofspiel_c4	num_cards=4, imp_info=True, 2 players, descending points	Zero-sum
first_sealed_auction_mv5	max_value=5, 2 players	General-sum
sheriff_small	max_items=2, max_bribe=2, num_rounds=2	General-sum

Table 1. Customized smaller games used in our experiments.

Game	Key parameters
kuhn_poker	2 players, 3-card deck (J/Q/K), 1 private card per player, no public card, 1 betting round
leduc_poker	2 players, 6-card deck (3 ranks $\times$ 2 suits), 1 private card per player, 1 public card, 2 betting rounds
universal_poker	2 players, 4 ranks, 2 suits, 2 rounds, 1 private card per player, public cards = (0, 1)

Table 2. Poker-family tasks and their explicit parameterizations.

Game	Key parameters	Utility
liars_dice	2 players, 1 six-sided die per player	General-sum

Table 3. Other standard imperfect-information games used in our experiments.

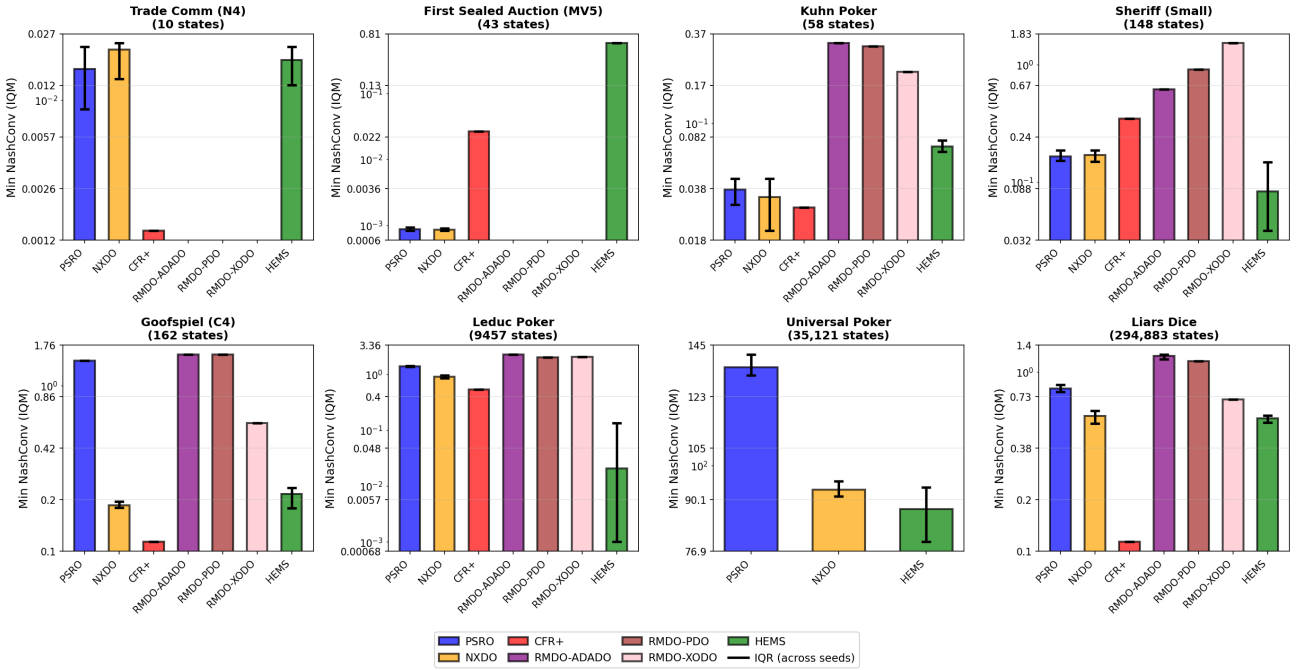

Figure 1. Per-task comparison of algorithms using the interquartile mean (IQM) of the *minimum* NashConv (exploitability) achieved per random seed. Each subplot corresponds to one game; tasks are ordered by increasing reported state-space size (annotated in the subplot title). Bars show IQM; vertical error bars indicate the interquartile range (IQR) across seeds. Algorithms include PSRO, NXDO, CFR$^+$, three RMDO variants (ADADO, PDO, XODO), and HEMS. The y-axis uses a logarithmic scale. One sentence: For Universal Poker we omit CFR$^+$ and the RMDO variants because the current implementation does not support running those baselines on this game, so no results were produced for those methods.

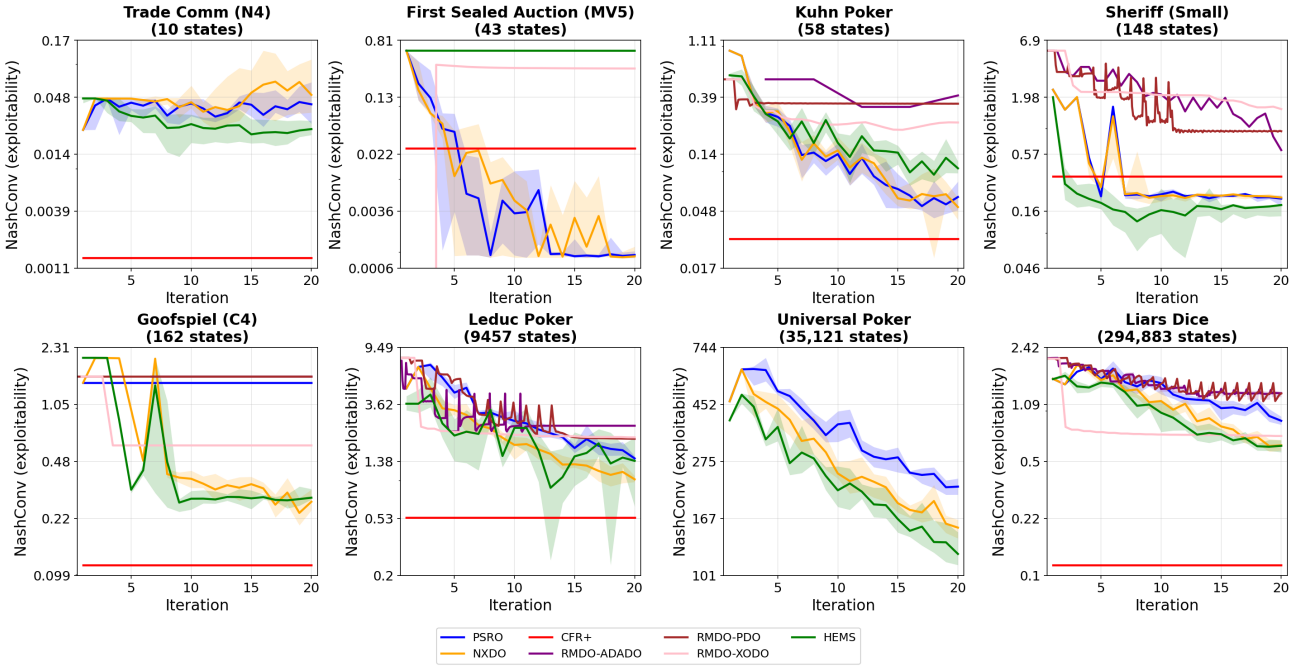

Figure 2. Exploitability (NashConv) as a function of training iteration for the same set of games as Fig. 1, in the same left-to-right order. Curves report the IQM across seeds at each iteration, with a shaded band for the IQR. For PSRO, NXDO, and HEMS, the first 20 iterations are shown on the horizontal axis. For CFR$^+$, a horizontal baseline denotes the final NashConv per seed, summarized by IQM with an IQR band. For the three RMDO variants, all stored iterations are used but the iteration index is linearly rescaled to $[1, 20]$ for comparability. The y-axis uses a logarithmic scale.